

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

导入数据

In [2]:


```
f= open("机器学习考核数据集.csv")
df= pd.read_csv(f)
print (df.shape)
df.head()
```

(41188, 21)

Out[2]:

	age	job	marital	education	default	housing	loan	contact	month	day_of
0	44	blue-collar	married	basic.4y	unknown	yes	no	cellular	aug	
1	53	technician	married	unknown	no	no	no	cellular	nov	
2	28	management	single	university.degree	no	yes	no	cellular	jun	
3	39	services	married	high.school	no	no	no	cellular	apr	
4	55	retired	married	basic.4y	no	yes	no	cellular	aug	

5 rows × 21 columns



步骤一：数据概览

In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
age                41188 non-null int64
job                41188 non-null object
marital            41188 non-null object
education          41188 non-null object
default            41188 non-null object
housing            41188 non-null object
loan               41188 non-null object
contact            41188 non-null object
month              41188 non-null object
day_of_week        41188 non-null object
duration           41188 non-null int64
campaign           41188 non-null int64
pdays             41188 non-null int64
previous           41188 non-null int64
poutcome           41188 non-null object
emp_var_rate       41188 non-null float64
cons_price_idx     41188 non-null float64
cons_conf_idx      41188 non-null float64
euribor3m          41188 non-null float64
nr_employed        41188 non-null float64
y                  41188 non-null int64
dtypes: float64(5), int64(6), object(10)
memory usage: 6.6+ MB
```

In [4]:

```
df.describe()
```

Out[4]:

	age	duration	campaign	pdays	previous	emp_var_rate	con
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	0.000000
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.000000
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	0.000000
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	0.000000
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	0.000000
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	0.000000
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	0.000000

查看各列值的分布

In [5]:

```
for column in df.columns:
    print("列名: %s"%df[column].name)
    print(df[column].value_counts())
    print("——"*15)
```

```
90      170
136      168
73       167
124      164
87       162
72       161
104      161
111      160
106      159
109      158
97       158
122      157
135      156
92       156
114      156
139      155
96       155
119      155
82       154
89       153
```

步骤二：缺失值处理

虽然在df.info()中看到数据都是non-null值，但在df[column].value_counts()中看到一些列有unknown值，以及“pdays”列有很多“999”的值，均表示缺失值，需要处理。

查看数据中“unknown”值个数

In [6]:

```
for column in df.columns:
    count=0
    for i in df[column]:
        if i=="unknown":
            count=count+1
    if count>0:
        print("列: %s, 其unknown值有%d个"%(df[column].name, count))
```

```
列: job, 其unknown值有330个
列: marital, 其unknown值有80个
列: education, 其unknown值有1731个
列: default, 其unknown值有8597个
列: housing, 其unknown值有990个
列: loan, 其unknown值有990个
```

看到有六列：job、marital、education、default（是否有违约）、housing（是否有房贷）、loan（是否有个人贷款）有unknown值，需处理

“default”列有8597个缺失值，显然，将含缺失值的行删除是不现实的，需想办法填充。

查看这些含“unknown”值的列的分布

In [7]:

```
l1=['job','marital','education','default','housing','loan']
for column in l1:
    print("列名: %s"%df[column].name)
    print(df[column].value_counts())
    print("——"*15)
```

```
列名: job
admin.          10422
blue-collar     9254
technician      6743
services        3969
management      2924
retired         1720
entrepreneur    1456
self-employed   1421
housemaid       1060
unemployed      1014
student         875
unknown         330
Name: job, dtype: int64
```

```
列名: marital
married         24928
single          11568
divorced         4612
unknown          80
Name: marital, dtype: int64
```

```
列名: education
university.degree  12168
high.school        9515
basic.9y           6045
professional.course  5243
basic.4y           4176
basic.6y           2292
unknown            1731
illiterate          18
Name: education, dtype: int64
```

```
列名: default
no              32588
unknown         8597
yes              3
Name: default, dtype: int64
```

```
列名: housing
yes            21576
no             18622
unknown        990
Name: housing, dtype: int64
```

```
列名: loan
no            33950
yes           6248
unknown       990
Name: loan, dtype: int64
```

为方便处理缺失值，将“unknown”全部替换为np.nan

In [8]:

```
df=df.replace("unknown", np. nan)
```

In [9]:

```
df.head()
```

Out[9]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_w
0	44	blue-collar	married	basic.4y	NaN	yes	no	cellular	aug	
1	53	technician	married	NaN	no	no	no	cellular	nov	
2	28	management	single	university.degree	no	yes	no	cellular	jun	
3	39	services	married	high.school	no	no	no	cellular	apr	
4	55	retired	married	basic.4y	no	yes	no	cellular	aug	

5 rows × 21 columns

ffill方法：用后一个非缺失值去填充该缺失值

In [10]:

```
df=df.fillna(method=' bfill')
l1=[' job', 'marital', ' education', ' default', ' housing', ' loan', ' pdays']
for column in l1:
    print("列名: %s"%df[column].name)
    print(df[column].value_counts())
    print("——"*15)
```

```
列名: job
admin.          10512
blue-collar     9310
technician      6814
services        4002
management      2942
retired         1732
entrepreneur    1467
self-employed   1432
housemaid       1070
unemployed      1021
student         886
Name: job, dtype: int64
```

```
列名: marital
married        24982
single         11586
divorced        4620
Name: marital, dtype: int64
```

看到：此时已经没有“unknown”/“nan”值了。

步骤三：离散特征编码

In [11]:

```
dic1={'yes':1,"no":0}
df["weiyue"]=df['default'].map(dic1)
df["housing_map"]=df['housing'].map(dic1)
df["loan_map"]=df['loan'].map(dic1)
dic2={'cellular':1,"telephone":0}
df["contact_map"]=df['contact'].map(dic2)
dic3={'success':2,"failure":0,"nonexistent":1}
df["poutcome_map"]=df['poutcome'].map(dic3)
```

In [12]:

```
df["default"].value_counts()
```

Out[12]:

```
no      41185
yes         3
Name: default, dtype: int64
```

In [13]:

```
df["weiyue"].value_counts()
```

Out[13]:

```
0      41185
1         3
Name: weiyue, dtype: int64
```

In [14]:

```
# 去除原“default”、“housing”、“loan”、“contact”、“poutcome”列
df.drop(['default','housing','loan','contact','poutcome'],axis=1, inplace=True)
```

删除无用列

In [15]:

```
# “month”（最近一次联系的月份）、“day_of_week”（最近一次联系的星期）是不能反映有用信息的，故将其删除
df.drop(['month','day_of_week'],axis=1, inplace=True)
#因为“pdays”这一列含有39673个“999”值，缺失值占比达96.32%，故删除该列
df.drop('pdays',axis=1, inplace=True)
```

get_dummies离散特征编码

In [16]:

```
df1=pd.get_dummies(df,prefix=['job','marital','education'])
df1.columns
```

Out[16]:

```
Index(['age', 'duration', 'campaign', 'previous', 'emp_var_rate',
      'cons_price_idx', 'cons_conf_idx', 'euribor3m', 'nr_employed', 'y',
      'weiyue', 'housing_map', 'loan_map', 'contact_map', 'poutcome_map',
      'job_admin.', 'job_blue-collar', 'job_entrepreneur', 'job_housemaid',
      'job_management', 'job_retired', 'job_self-employed', 'job_services',
      'job_student', 'job_technician', 'job_unemployed', 'marital_divorced',
      'marital_married', 'marital_single', 'education_basic.4y',
      'education_basic.6y', 'education_basic.9y', 'education_high.school',
      'education_illiterate', 'education_professional.course',
      'education_university.degree'],
      dtype='object')
```

In [17]:

```
#np.isnan(df1).any()          #验证df1已没有缺失值
```

步骤四：连续值分组处理

4.1 “age”年龄分组

In [18]:

```
bins=[0, 20, 40, 60, 98]
labels=["20岁以下", "20到40岁", "40到60岁", "60岁以上"]
df1["age_group"]=pd.cut(df1.age, bins, labels=labels)
```

In [19]:

```
dic10={'20岁以下':1, "20到40岁":2, "40到60岁":3, "60岁以上":4}
df1["age_map"]=df1['age_group'].map(dic10)
```


In [20]:

```
df1.age_group
```

Out[20]:

0	40到60岁
1	40到60岁
2	20到40岁
3	20到40岁
4	40到60岁
5	20到40岁
6	20到40岁
7	20到40岁
8	20到40岁
9	20到40岁
10	20到40岁
11	40到60岁
12	40到60岁
13	20到40岁
14	20到40岁
15	40到60岁
16	20到40岁
17	20到40岁
18	20到40岁
19	20到40岁
20	40到60岁
21	20到40岁
22	20到40岁
23	20到40岁
24	40到60岁
25	20到40岁
26	40到60岁
27	20到40岁
28	20到40岁
29	20到40岁
...	
41158	20到40岁
41159	40到60岁
41160	40到60岁
41161	40到60岁
41162	40到60岁
41163	20到40岁
41164	20到40岁
41165	40到60岁
41166	20到40岁
41167	20到40岁
41168	20到40岁
41169	40到60岁
41170	20到40岁
41171	40到60岁
41172	20到40岁
41173	20到40岁
41174	40到60岁
41175	40到60岁
41176	20到40岁
41177	40到60岁
41178	20到40岁
41179	20到40岁
41180	40到60岁
41181	20到40岁
41182	20到40岁

```
41183    40到60岁
41184    20到40岁
41185    40到60岁
41186    40到60岁
41187    20到40岁
```

```
Name: age_group, Length: 41188, dtype: category
```

```
Categories (4, object): [20岁以下 < 20到40岁 < 40到60岁 < 60岁以上]
```

交叉分析：观察age_group与y的分布

In [21]:

```
df1.pivot_table(values=["age"], index=["age_group"], columns=["y"], aggfunc=[np.size])
```

Out[21]:

y	size	
	age	
	0	1
age_group		
20岁以下	83	57
20到40岁	20964	2664
40到60岁	15005	1505
60岁以上	496	414

4.2 “duration”通话持续时间分组

In [22]:

```
bins=[min(df1.duration)-1, 102, 180, 320, max(df1.duration)+1] #这里的102、180、320分别为df.duration
labels=["102分钟以下", "102到180分钟", "180到320分钟", "320分钟及以上"]
df1["duration_group"]=pd.cut(df1.duration, bins, labels=labels)
```

In [23]:

```
dic4={'102分钟以下':1, "102到180分钟":2, "180到320分钟":3, "320分钟及以上":4}
df1["duration_map"]=df1["duration_group"].map(dic4)
```

交叉分析：观察duration_group与y的分布

可大致观察到，y=1的类别里，通话分钟数越长，注册定期存款的人数越多。

In [24]:

```
df1.pivot_table(values=["duration"], index=["duration_group"], columns=["y"], aggfunc=[np.size])
```

Out[24]:

	size	
	duration	
y	0	1
duration_group		
102分钟以下	10226	87
102到180分钟	9915	477
180到320分钟	9163	1081
320分钟及以上	7244	2995

In [25]:

```
df1.drop('duration_group', axis=1, inplace=True)
df1.drop('duration', axis=1, inplace=True)
```

4.3 “campaign”联系次数分组

In [26]:

```
bins=[min(df1.campaign)-1, 3, 6, 10, max(df1.campaign)+1]
labels=["0-2次", "3-5次", "6-9次", "10次以上"]
df1["campaign_group"]=pd.cut(df1.campaign, bins, labels=labels, right=False)
```

In [27]:

```
dic5={"0-2次":1, "3-5次":2, "6-9次":3, "10次以上":4}
df1["campaign_map"]=df1['campaign_group'].map(dic5)
df1.drop('campaign_group', axis=1, inplace=True)
df1.drop('campaign', axis=1, inplace=True)
```

In [28]:

df1.head()

Out[28]:

	age	previous	emp_var_rate	cons_price_idx	cons_conf_idx	euribor3m	nr_employed	y	w
0	44	0	1.4	93.444	-36.1	4.963	5228.1	0	
1	53	0	-0.1	93.200	-42.0	4.021	5195.8	0	
2	28	2	-1.7	94.055	-39.8	0.729	4991.6	1	
3	39	0	-1.8	93.075	-47.1	1.405	5099.1	0	
4	55	1	-2.9	92.201	-31.4	0.869	5076.2	1	

5 rows × 38 columns

步骤五：数据标准化

In [29]:

```
X=df1[['previous', 'emp_var_rate',
        'cons_price_idx', 'cons_conf_idx', 'euribor3m', 'nr_employed',
        'weiyue', 'housing_map', 'loan_map', 'contact_map', 'poutcome_map',
        'job_admin.', 'job_blue-collar', 'job_entrepreneur', 'job_housemaid',
        'job_management', 'job_retired', 'job_self-employed', 'job_services',
        'job_student', 'job_technician', 'job_unemployed', 'marital_divorced',
        'marital_married', 'marital_single', 'education_basic.4y',
        'education_basic.6y', 'education_basic.9y', 'education_high.school',
        'education_illiterate', 'education_professional.course',
        'education_university.degree']]
```

In [30]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled
```

Out[30]:

```
array([[ -0.34949428,  0.83906065, -0.2274652 , ..., -0.02204119,
        -0.39125376, -0.66909041],
       [ -0.34949428, -0.11578127, -0.649003  , ..., -0.02204119,
        -0.39125376,  1.49456633],
       [  3.69176641, -1.13427931,  0.82810692, ..., -0.02204119,
        -0.39125376,  1.49456633],
       ...,
       [ -0.34949428,  0.64809227,  0.72272247, ..., -0.02204119,
        -0.39125376,  1.49456633],
       [ -0.34949428, -2.21643348, -1.97753812, ..., -0.02204119,
         2.55588596, -0.66909041],
       [ -0.34949428,  0.64809227,  0.72272247, ..., -0.02204119,
        -0.39125376, -0.66909041]])
```

步骤六：逻辑回归预测

模型训练

In [31]:

```
from sklearn.linear_model import LogisticRegression
lr1=LogisticRegression()
lr1.fit(X, df1["y"])
```

Out[31]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)
```

模型预测 (predict_proba方法)

In [32]:

```
pred=lr1.predict(X)
print(pred)
```

```
[0 0 1 ..., 0 0 0]
```

模型评判之——精度值计算

In [33]:

```
df1["pred_label"]=pred
```

In [34]:

```
match=df1["y"]==df1["pred_label"]
correct=df1[match]
accuracy=len(correct)/len(df1)
print(accuracy)
```

```
0.8989754297368165
```

模型评判之——TPR计算

In [35]:

```
TP=(df1["y"]==1)&(df1["pred_label"]==1)
TP=len(df1[TP])
print(TP)
FN=(df1["y"]==1)&(df1["pred_label"]==0)
FN=len(df1[FN])
print(FN)
TPR=TP/(TP+FN)
print(TPR)
```

```
860
3780
0.1853448275862069
```

模型评判之——TNR计算

In [36]:

```
TN=(df1["y"]==0)&(df1["pred_label"]==0)
TN=len(df1[TN])
print(TN)
FP=(df1["y"]==0)&(df1["pred_label"]==1)
FP=len(df1[FP])
print(FP)
TNR=TN/(TN+FP)
print(TNR)
```

```
36167
381
0.9895753529604903
```

ROC指标 (metrics.roc_curve)

In [37]:


```
from sklearn import metrics

false_positive_rate,true_positive_rate,thresholds =metrics.roc_curve(df1["y"],df1["pred_label"])
print(thresholds)
plt.plot(false_positive_rate,true_positive_rate)
```

```
[2 1 0]
```

Out[37]:

[<matplotlib.lines.Line2D at 0xe201e10>]



计算AUC值 (即ROC面积) (sklearn.metrics.roc_auc_score)

In [38]:

```
from sklearn.metrics import roc_auc_score

auc_score=roc_auc_score(df1["y"],df1["pred_label"])
print(auc_score)
```

```
0.587460090273
```

交叉验证 (sklearn.cross_validation)

In [39]:

```
from sklearn.cross_validation import KFold
from sklearn.cross_validation import cross_val_score

lr=LogisticRegression()
kf=KFold(len(df1), 5, shuffle=True)
print("每次交叉验证的准确率为: ")
accuracies=cross_val_score(lr, X, df["y"], scoring="accuracy", cv=kf)
print(accuracies)
print(type(accuracies))
accuracy=accuracies.mean()
print("平均准确率为: ")
print(accuracy)
```

每次交叉验证的准确率为:

D:\anzhuangruanjian\anaconda\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

[0.89900461 0.90082544 0.89839767 0.89802112 0.89328639]

<class 'numpy.ndarray'>

平均准确率为:

0.89790704801

In []:

In []:

In []:

In []:

In []:

In []:

In []: